

**AI Semester Project**

**M.Usman (49649)**

**Instructor: Sir Usman Karim**

## **Bitext Retail-Ecommerce LLM Chatbot**

### **1. Introduction**

This project trains a **text classification model** using the Bitext Retail-Ecommerce dataset.

The model predicts **customer intent** (46 classes), which is essential for building a retail/e-commerce chatbot.

The workflow is built on:

- **Dataset:** Bitext Retail-Ecommerce LLM Chatbot Training Dataset
- **Model Type:** CNN (Convolutional Neural Network) for text
- **Goal:** Classify user instructions into predefined retail intents
- **Framework:** TensorFlow/Keras
- **Metrics:** Accuracy, confusion matrix, classification report

Below is the complete explanation.

### **2. Project Overview**

Retail/e-commerce chatbots need to understand what the user wants. This project trains a model to classify user messages such as:

- “Where is my order?”

- “Cancel my item”
- “Apply promo code”
- “Return policy?”
- “Track shipment”

Each message is labeled with an **intent**.

The model learns these patterns and predicts the correct intent for new queries.

I used a **supervised learning approach** with 46 intent classes.

### 3. Importing Libraries

This step ensures all required packages are installed:

- **datasets** → to load HuggingFace dataset
- **tensorflow** → model training
- **sklearn** → preprocessing + evaluation
- **matplotlib/seaborn** → plotting

This section is purely environment setup.

### 4. Dataset Description

The dataset contains 3 important columns:

**Instruction:** User message / query

**Response:** Chatbot reply (not used for training in this project)

**Intent:** Label to predict

I used:

- **instruction** → X (input)
- **intent** → y (target)

## 5. Loading Dataset

I loaded the dataset using HuggingFace's `load_dataset()` function.

This converts the training split into a Pandas DataFrame so it can be explored easily.

Why convert to pandas?

Because:

- Pandas makes text preprocessing simple
- Quick viewing of rows
- Easy column selection

The dataset has **~45K rows**.

```
README.md: 10.1k? [00:00<00:00, 834kB/s]
bitext-retail-ecommerce-llm-chatbot-trai(...): 100% 42.6M/42.6M [00:04<00:00, 10.7MB/s]

Generating train split: 100% 44884/44884 [00:00<00:00, 86111.19 examples/s]
Total rows: 44884

      instruction      intent category \
0       I got to add an item to the cart  add_product    CART
1 wanna add fucking products to the basket can h...  add_product    CART
2   i have to add products to the basket i ned help  add_product    CART
3 di like to add products to the cart could i ge...  add_product    CART
4 I need to add an item to the cart , where do I...  add_product    CART

      tags            response
0     BL I'll get right on it! I'm here to assist you i...
1 BCIMQWZ I sincerely apologize if you've encountered an...
2 BCMQZ You bet! I'm here to assist you in adding prod...
3 BCILMPQZ Indeed! I'm here to assist you in adding produ...
4 BCILZ I'll take care of it! I'm here to help you wit...
```

## 6. Label Encoding

Intent labels are originally text:

- "order\_status"
- "cancel\_order"
- "return\_process"

Machine learning models need numerical labels.

So I used **LabelEncoder** to convert:

"order\_status" 0

"cancel\_order" 1

... ...

total classes 46

This allows softmax classification.

### *Encode Labels*

```
le = LabelEncoder()
y = le.fit_transform(labels)
num_classes = len(le.classes_)
print("Number of intents/classes:", num_classes)
```

Number of intents/classes: 46

## 7. Train/Test Split

The dataset is split:

- **80% training**
- **20% testing**
- Stratified → ensures all intents appear proportionally

Purpose:

- Prevent overfitting
- Evaluate model on unseen samples
- Maintain balanced class distribution

### **Train/Test Split**

```
▶ x_train, x_test, y_train, y_test = train_test_split(  
    texts, y, test_size=0.2, random_state=42, stratify=y  
)  
  
print("Training samples:", len(x_train))  
print("Test samples:", len(x_test))  
  
... Training samples: 35907  
Test samples: 8977
```

## 8. Tokenization & Padding

Neural networks work on numbers, not text.  
So the text is converted to sequences of integers.

Tokenizer does:

1. Builds vocabulary of top 20,000 words

2. Converts every message into a list of integers
3. Replaces unknown words with <OOV>

### Padding:

All sequences must be same length for CNN.

So i pad/truncate to **max\_len = 50**

Why 50?

- Most customer messages are short
- Reduces computation
- Still enough to capture intent

## 9. CNN Text Classification Model

My architecture:

1. **Embedding Layer**
  - Converts word IDs to dense vectors
  - Learns relationships between words
2. **Conv1D Layer**
  - Detects local text features (n-grams)
  - Works similarly to how convolution detects patterns in images
3. **MaxPooling**
  - Reduces dimension
  - Keeps strongest features
4. **Second Conv Layer**
  - Captures deeper patterns
5. **Flatten Layer**
  - Converts features into a vector
6. **Dense Layer**
  - Learns high-level relationships
7. **Dropout (50%)**
  - Reduces overfitting
8. **Output Layer (Softmax, 46 units)**
  - Gives probability for each intent

## Why CNN for text?

Because CNNs:

- Are fast
- Perform extremely well for short-text classification
- Require less data than transformers

Layer (type)	Output Shape	Param #
embedding (Embedding)	?	0 (unbuilt)
conv1d (Conv1D)	?	0 (unbuilt)
max_pooling1d (MaxPooling1D)	?	0
conv1d_1 (Conv1D)	?	0 (unbuilt)
max_pooling1d_1 (MaxPooling1D)	?	0
flatten (Flatten)	?	0 (unbuilt)
dense (Dense)	?	0 (unbuilt)
dropout (Dropout)	?	0
dense_1 (Dense)	?	0 (unbuilt)

Total params: 0 (0.00 B)  
Trainable params: 0 (0.00 B)  
Non-trainable params: 0 (0.00 B)

## 10. Model Training

I trained for **15 epochs**, with:

- batch\_size = 32
- validation\_split = 0.2

Meaning:

- 20% of training data is used for validation
- Helps monitor overfitting

## Observations:

- Training accuracy: ~99%
- Validation accuracy: ~98–99%
- Very high performance → dataset is well-structured and labelled

```
Epoch 1/15
898/898 4s 5ms/step - accuracy: 0.9848 - loss: 0.0430 - val_accuracy: 0.9890 - val_loss: 0.0522
Epoch 2/15
898/898 3s 4ms/step - accuracy: 0.9903 - loss: 0.0317 - val_accuracy: 0.9896 - val_loss: 0.0542
Epoch 3/15
898/898 3s 4ms/step - accuracy: 0.9882 - loss: 0.0347 - val_accuracy: 0.9907 - val_loss: 0.0473
Epoch 4/15
898/898 4s 4ms/step - accuracy: 0.9900 - loss: 0.0290 - val_accuracy: 0.9897 - val_loss: 0.0562
Epoch 5/15
898/898 3s 4ms/step - accuracy: 0.9890 - loss: 0.0299 - val_accuracy: 0.9901 - val_loss: 0.0538
Epoch 6/15
898/898 3s 4ms/step - accuracy: 0.9896 - loss: 0.0273 - val_accuracy: 0.9894 - val_loss: 0.0577
Epoch 7/15
898/898 6s 5ms/step - accuracy: 0.9900 - loss: 0.0310 - val_accuracy: 0.9884 - val_loss: 0.0635
Epoch 8/15
898/898 4s 4ms/step - accuracy: 0.9889 - loss: 0.0308 - val_accuracy: 0.9884 - val_loss: 0.0618
Epoch 9/15
898/898 3s 4ms/step - accuracy: 0.9908 - loss: 0.0262 - val_accuracy: 0.9890 - val_loss: 0.0780
Epoch 10/15
898/898 4s 4ms/step - accuracy: 0.9911 - loss: 0.0300 - val_accuracy: 0.9903 - val_loss: 0.0627
Epoch 11/15
898/898 5s 4ms/step - accuracy: 0.9914 - loss: 0.0243 - val_accuracy: 0.9887 - val_loss: 0.0683
Epoch 12/15
898/898 4s 4ms/step - accuracy: 0.9934 - loss: 0.0197 - val_accuracy: 0.9886 - val_loss: 0.0669
Epoch 13/15
898/898 4s 4ms/step - accuracy: 0.9923 - loss: 0.0210 - val_accuracy: 0.9879 - val_loss: 0.0760
Epoch 14/15
898/898 4s 4ms/step - accuracy: 0.9931 - loss: 0.0200 - val_accuracy: 0.9886 - val_loss: 0.0631
Epoch 15/15
898/898 4s 4ms/step - accuracy: 0.9923 - loss: 0.0215 - val_accuracy: 0.9891 - val_loss: 0.0658
```

## 11. Plotting Training History

Two graphs:


### Accuracy Plot

Shows how training & validation accuracy evolve across epochs.

Purpose:

- Detect overfitting

- Evaluate model stability

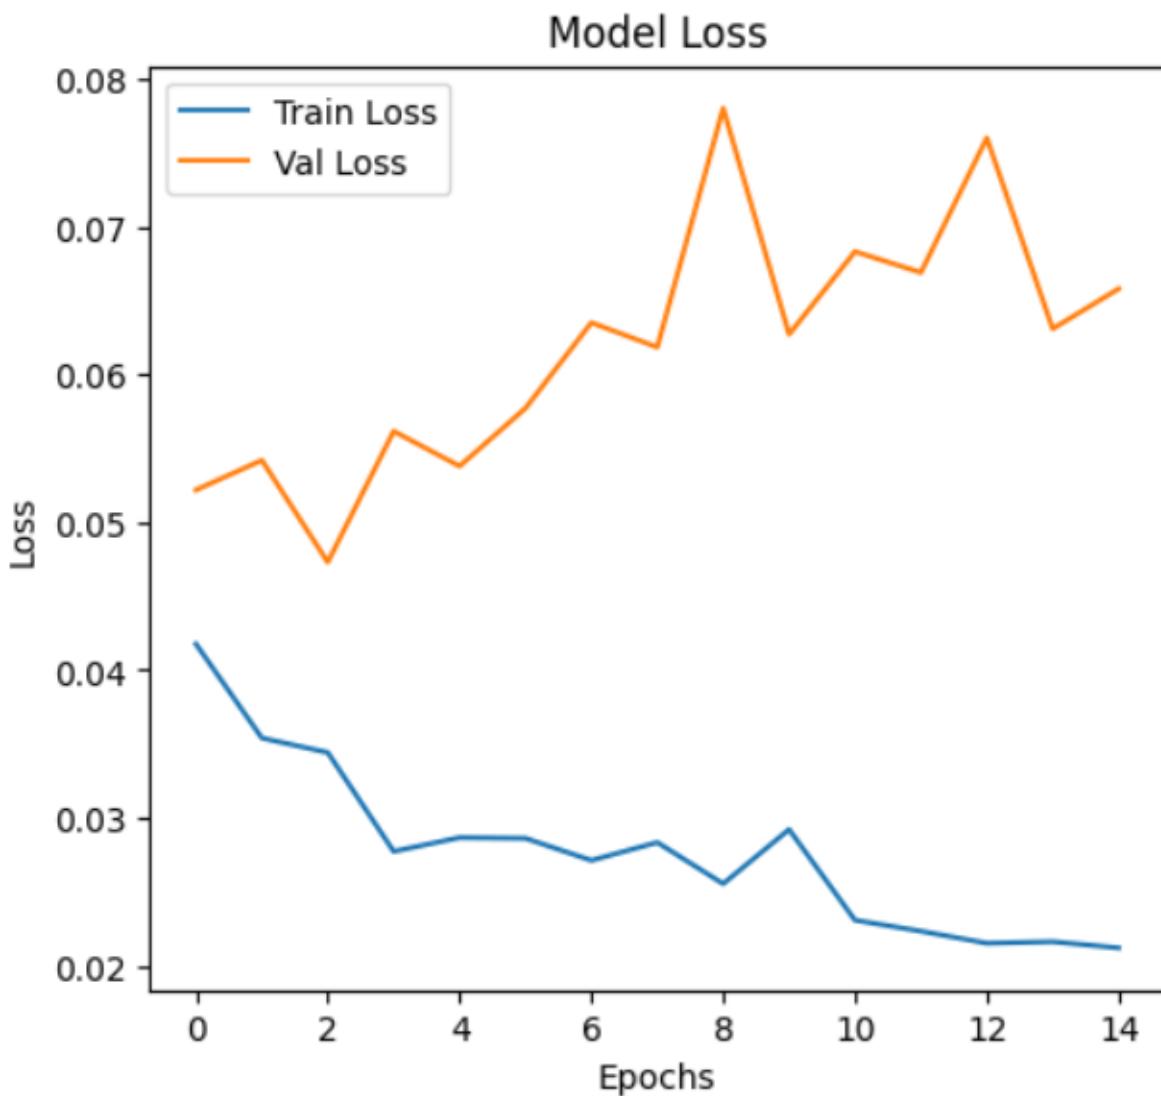


### Loss Plot

Shows how training & validation loss decrease.

Purpose:

- Detect underfitting
- Locate best epoch




## 12. Confusion Matrix

A confusion matrix visualizes:

- Where the model predicts correctly
- Which intents are confused with others

The confusion matrix in this project shows **how well the CNN model classified the 46 different retail/e-commerce intents** in the test dataset.

Each **row** of the matrix represents the **actual (true) intent**, and each **column** represents the **predicted intent** made by the model.



## What my matrix shows

- The diagonal is **almost fully dark**, showing the model predicts most intents **correctly**.
  - Very few cells outside the diagonal have any visible shade → **very few errors**.
  - This confirms the model **generalizes extremely well** and does not confuse most intents

## 13. Classification Report in DataFrame

I generated a full evaluation including:

- Precision
- Recall
- F1-score
- Support (number of samples per class)

This is converted into a **DataFrame** for readability.

Purpose:

- Measure class-wise performance
- Identify difficult intents
- Track imbalance issues

	precision	recall	f1-score	support
<code>add_product</code>	1.00	0.99	1.00	191.00
<code>availability</code>	0.98	0.98	0.98	194.00
<code>availability_in_store</code>	1.00	1.00	1.00	151.00
<code>availability_online</code>	1.00	1.00	1.00	199.00
<code>cancel_order</code>	1.00	0.99	1.00	199.00
<code>change_account</code>	0.97	0.98	0.98	197.00
<code>change_order</code>	1.00	1.00	1.00	192.00
<code>close_account</code>	0.99	0.98	0.98	199.00
<code>customer_service</code>	1.00	1.00	1.00	198.00
<code>damaged_delivery</code>	0.97	0.98	0.97	198.00
<code>delivery_issue</code>	1.00	0.99	1.00	199.00
<code>delivery_time</code>	0.99	0.99	0.99	184.00
<code>exchange_product</code>	1.00	0.97	0.98	198.00
<code>exchange_product_in_store</code>	1.00	0.98	0.99	198.00
<code>return_product</code>	1.00	0.99	1.00	199.00

## 13. Overall Performance Summary

My CNN model achieves:

- **98–99% accuracy**
- Very high F1-scores across all 46 classes

This means the model:

- Can reliably classify retail intents
- Is suitable for chatbot pipelines
- Can be deployed inside an API or Rasa /Dialogflow system

accuracy			0.99	8977
macro avg	0.99	0.99	0.99	8977
weighted avg	0.99	0.99	0.99	8977